

UNCERTAINTY IN RECURRENT NEURAL NETWORKS

ALIREZA SAMAR

A thesis submitted in fulfilment of the
requirements for the award of the degree of
Master of Philosophy

Advanced Informatics School
Universiti Teknologi Malaysia

APRIL 2017

ABSTRACT

Recurrent neural networks (RNN), and especially Deep RNNs have outperformed in various fields from computer vision, and language processing to physics, biology, and manufacturing. This means the deep or multi-layer architecture of neural networks are being extensively used in these fields; for instance convolutional neural networks (CNN) as image processing tools, and recurrent neural networks (RNN) as sequence processing model various from language modeling to image captioning.

However, in traditional sciences fields such as physics and biology, analysis model uncertainty is crucial, especially in time series models where delay cant be tolerated. This observation and probabilistic models provides the confidence bounds for understanding data and making the decision based on analysis.

This work aims to propose a novel theoretical framework and develop tools to measure uncertainty estimates by adoption of Bayes by Backprop to the given network, especially in deep recurrent neural networks. The performance of proposed technique in this work will be extensively evaluated with widely studied benchmark.

TABLE OF CONTENTS

CHAPTER	TITLE	PAGE
	ABSTRACT	3
	TABLE OF CONTENTS	5
	LIST OF TABLES	7
	LIST OF FIGURES	9
	LIST OF ABBREVIATIONS	11
	LIST OF SYMBOLS	13
1	INTRODUCTION	1
	1.1 Introduction: The Importance of Uncertainty	1
	1.2 Problem Background	1
	1.3 Problem Statement	2
	1.4 Project Aim	2
	1.5 Project Questions	2
	1.6 Objective and Scope	3
2	LITERATURE REVIEW	5
	2.1 Introduction	5
	2.1.1 Recurrent Neural Networks	5
	2.2 Similar Works	5
	2.3 State-of-the-Arts	6
	2.4 Limitations	7
	2.5 Research Gaps	7
3	RESEARCH METHODOLOGY	9
	3.1 Technique: Backprop Through Time	9
	3.2 Benchmarks	10
	3.2.1 Language Modeling	10

4	ANALYSIS AND DESIGN	13
4.1	Bayes by Backprop	13
4.2	Adoption of Bayes by Backprop to RNNs	14
	REFERENCES	17

LIST OF TABLES

TABLE NO.	TITLE	PAGE
3.1	Word-level perplexity on the Penn Treebank language modelling task (lower is better).	11

LIST OF FIGURES

FIGURE NO.	TITLE	PAGE
2.1	A Rolled Recurrent Neural Networks	8
2.2	An Unrolled Recurrent Neural Networks	8
4.1	Adoption of Bayes by Backprop (BBB) to an RNN.	15

LIST OF ABBREVIATIONS

ANN	-	Artificial Neural Network
NN	-	Neural Network
BNN	-	Bayesian Neural Network
RNN	-	Recurrent Neural Network
BBB	-	Bayes by Backprop
CNN	-	Convolutional Neural Network
KL	-	Kullback-Leibler
GP	-	Gaussian Process
MC	-	Monte Carlo
VI	-	Variational Inference
MCMC	-	Markov Chain Monte Carlo
LSTM	-	Long Short-Term Memory
SRT	-	Stochastic Regularisation Technique (such as dropout)
i.e.	-	Id est (it is)
w.r.t.	-	With respect to

LIST OF SYMBOLS

A	-	A rolled NN
\mathbf{A}	-	Matrix
\mathbf{a}	-	vector
t	-	Steps/Length of RNN
x_t	-	Input value
η	-	Trained parameter
c	-	Internal core sate
h	-	Exposed sate
i_t	-	Input gate
f_t	-	Forget gate
h_t	-	Exposed sate
W	-	Weights (biases)

CHAPTER 1

INTRODUCTION

1.1 Introduction: The Importance of Uncertainty

The Bayesian approach to machine learning is based on using probability to represent all forms of uncertainty. There are different models like the Gaussian process to understand possible likely and less likely options to generalize the observed data by defining the probability of distributions over functions. This observation and probabilistic models provides the confidence bounds for understanding data and making the decision based on analysis. For instance, an autonomous vehicle would use the determination from confidence bounds to whether brake or not. The confidence bounds simply means *how certain the model is about its output?*

Understanding whether the chosen model is the right one or not, or the data has enough signals or not is an active field of research [1] in *Bayesian machine learning*, especially in *deep learning models* where based on predictions result it's difficult to make sure about the certainty level of predictions.

1.2 Problem Background

Recurrent Neural Networks (RNNs) achieve state-of-the-art performance on a wide range of sequence prediction tasks ([2]; [3]; [4]; [5]; [6]). In this work we shall examine how to add uncertainty and regularisation to RNNs by means of applying Bayesian methods to training. Bayesian methods give RNNs another way to express their uncertainty (via the parameters). At the same time, by using a prior to integrate out the parameters to average across many models during training, this gives a regularisation effect to the network. Recent approaches either attempt to justify

dropout [7] and weight decay as a variational inference scheme [8], or apply Stochastic Gradient Langevin dynamics [9] to truncated backpropagation in time directly [10].

1.3 Problem Statement

Interestingly, recent work has not explored further directly apply a variational Bayes inference scheme for RNNs as was done in practical. We derive a straightforward approach based upon Bayes by Backprop [11] that we show works well on large scale problems. Our approach is a simple alteration to truncated backpropagation through time that results in an estimate of the posterior distribution on the weights of the RNN. Applying Bayesian methods to successful deep learning models affords two advantages: explicit representations of uncertainty and regularisation. Our formulation explicitly leads to a cost function with an information theoretic justification by means of a bits-back argument [12] where a KL divergence acts as a regulariser.

1.4 Project Aim

The aim of this work is to propose a novel theoretical framework and develop tools to measure uncertainty estimates by adoption of Bayes by Backprop to the given network, especially in deep recurrent neural networks and test the performance of proposed technique in this work with widely studied benchmarks.

1.5 Project Questions

The contributions of this work are as follows:

1. How to do uncertainty analysis in RNNs?
2. How to reduce the variance of probabilistic views?
3. How regularisation techniques can improve the performance?

1.6 Objective and Scope

The contributions of this work are as follows:

1. To investigate and demonstrate how Bayes by Backprop (BBB) can be efficiently applied to Recurrent Neural Networks (RNNs).
2. Develop a novel technique to reduces the variance of Bayes by Backprop (BBB).
3. Improve the performance on a widely studied benchmark with established regularisation technique such as dropout.

CHAPTER 2

LITERATURE REVIEW

2.1 Introduction

2.1.1 Recurrent Neural Networks

Recurrent Neural Networks (RNNs) are in forefront of recent development and advances in *deep learning* by making able neural networks to deal with sequences data, which is a major shortcoming in ANN. If the data is based on sequence of events in a video or text, the traditional neural network can't do reasoning for a single event based on its previous one. To tackle this issue RNNs have loops which enables them to persist the information.

As it shown in **Figure 2.1**, a selected neural network, A takes the input x_t and outputs the value of h_t . this might not show how data goes from one step to the next one in a same network until you unroll the loop and see chain architecture of recurrent neural networks that makes them the best choice for sequential data, **Figure 2.2**.

While RNNs is being used in variety of applications from language modeling to image captioning, the essential to all these achievement is the RNN-LSTMs [?]. An enhanced version of RNNs that outperforms better than the standard RNN.

2.2 Similar Works

Sharpening the posterior as proposed in Chapter 3 has mutualities with other techniques. Such as line search, where η is a trained parameter that does line search along the gradient direction to give probabilistic interpretations to line search. In this

work the novel technique is to use a variational posterior with the reparametrization gradient.

Another similar technique is dynamic evaluation [?], which trains the RNN during evaluation steps of the model with a fixed learning rate. The drawback of this technique is that each update applied in this case is cumulative, and only uses previously seen data.

Finally, learning to optimise [?] is considered as similar work also since it is learned so supposedly it produces better updates than the rest suggested by AdaGrad [?] or Adam [?]. Whilst they train a parametric model, we treat these as free parameters (so that they can adapt more quickly to the non-stationary distribution w.r.t. parameters). Notably, we use gradient information to inform a variational posterior so as to reduce variance of Bayesian Neural Networks. Thus, although similar in flavour, the underlying motivations are quite different.

2.3 State-of-the-Arts

Applying Bayesian methods to neural networks has a long history, with most common approximations having been tried. [?] propose various maximum a posteriori schemes for neural networks, including an approximate posterior centered at the mode. [?] also suggest using second order derivatives in the prior to encourage smoothness of the resulting network. [12] proposed using variational methods for compressing the weights of neural networks as a regulariser. [?] suggest an MDL loss for single layer networks that penalises non-robust weights by means of an approximate penalty based upon perturbations of the weights on the outputs. [?] investigated using the Laplace approximation for capturing the posterior of neural networks. [?] investigated the use of hybrid Monte Carlo for training neural networks, although it has so far been difficult to apply these to the large sizes of networks considered here.

More recently [?] derived a variational inference scheme for neural networks and [11] extended this with an update for the variance that is unbiased and simpler to compute. [?] derives a similar algorithm in the case of a mixture posterior. Several authors have claimed that dropout [7] and Gaussian dropout [?] can be viewed as approximate variational inference schemes [8], [?]. [10] goes a step further and uses Monte Carlo dropout for LSTMs (we compare to this results in our experiments).

Variational methods typically underestimate the uncertainty in the posterior (as they are mode seeking, akin to the Laplace approximation), whereas expectation propagation methods are mode averaging and so tend to overestimate uncertainty. Nonetheless, several papers explore applying expectation propagation to neural networks: [?] derive a closed form approximate online expectation propagation algorithm, whereas hernandez2015probabilistic proposed using multiple passes of assumed density filtering (in combination with early stopping) attaining good performance on a number of small data sets. hasenclever2015distributed derive a distributed expectation propagation scheme with SGLD [9] as an inner loop. Others have also considered applying SGLD to neural networks li2015preconditioned and [10] more recently used SGLD for LSTMs (we compare to these results in our experiments).

2.4 Limitations

1. What stage to sample the parameters of RNNs?
2. How to weight the contribution of KL regulariser?

2.5 Research Gaps

The quality of uncertainty estimates and the performance of the model in general is being affected by the form of posterior in variational inference. That being said, by adopting the posterior locally, or sharpening, to a batch we can improve the performance using gradients upon the batch.

This technique could give flexible form to the typical assumption of Gaussian posterior (GP), that variational inference is applied to neural networks to reduce variance.




Figure 2.1: Recurrent Neural Networks (RNNs) uses loops.




Figure 2.2: An Unrolled Recurrent Neural Networks (RNNs).

CHAPTER 3

RESEARCH METHODOLOGY

3.1 Technique: Backprop Through Time

As illustrated in 2.1, the core of an recurrent neural networks (RNNs), f , is a neural network that maps the RNN state at step t , s_t and an input observation x_t to a new RNN state s_{t+1} , $f : (s_t, x_t) \mapsto s_{t+1}$.

For comparison, an LSTM-powered RNN core [?] has a state $s_t = (c_t, h_t)$ where c is an internal core state and h is the exposed state. Intermediate gates modulate the effect of the inputs on the outputs, gates like the input gate i_t , forget gate f_t and output gate o_t . The correlation between the inputs, outputs and internal gates of an LSTM cell can be explained as:

$$\begin{aligned} i_t &= \sigma(W_i[x_t, h_{t-1}]^T + b_i), \\ f_t &= \sigma(W_f[x_t, h_{t-1}]^T + b_f), \\ c_t &= f_t c_{t-1} + i_t \tanh(W_c[x_t, h_{t-1}] + b_c), \\ o_t &= \sigma(W_o[x_t, h_{t-1}]^T + b_o), \\ h_t &= o_t \tanh(c_t), \end{aligned}$$

In the above statement, W_i (b_i), W_f (b_f), W_c (b_c) and W_o (b_o) are the weights (biases) that are affecting the input gate, forget gate, cell update, and output gate accordingly.

An RNN can be trained on a sequence of T using backpropagation through

time where the RNN is unrolled T times like a feed-forward network. Which can be achieved with forming the feed-forward network with inputs x_1, x_2, \dots, x_T and initial state s_0 :

$$\begin{aligned} s_1 &= f(s_0, x_1), \\ s_2 &= f(s_1, x_2), \\ &\dots \\ s_T &= f(s_{T-1}, x_T), \end{aligned} \tag{3.1}$$

where s_T is the total length (final state) of the RNN. Referring to the unrolled RNN for T steps as in (3.1) by $s_{1:T} = F_T(x_{1:T}, s_0)$, where $x_{1:T}$ is the sequence of input vectors and $s_{1:T}$ is the sequence of corresponding states. However, the truncated version of the algorithm can be seen as taking s_0 as the last state of the previous batch, s_T .

This RNN parameters are trained in the same way as a feed-forward neural network and a loss is applied to the states $s_{1:T}$ of the RNN, and then backpropagation is being used to update the weights of the trained network. Since the weights in each of the unrolled step are shared, each weight of the RNN core receives T gradient contributions when it is unrolled for T steps.

3.2 Benchmarks

3.2.1 Language Modeling

The experiments result of this work will be tested in Chapter 5, on the Penn Treebank [?] benchmark, a task consisting on next word prediction. Using the network architecture from [5] as a baseline and for which there is an open source implementation¹.

¹https://github.com/tensorflow/models/blob/master/tutorials/rnn/ptb/ptb_word_lm.py

Table 3.1: Word-level perplexity on the Penn Treebank language modelling task (lower is better).

Model (medium)	Validation	Test
LSTM	120.7	114.5
LSTM dropout	86.2	82.1
SGLD	N/A	109.1
SGLD dropout	N/A	99.6
Variational LSTM (tied weights)	81.8	79.7
Variational LSTM (tied weights, MS)	N/A	79.0
Model (medium) with Dynamic Evaluation	Validation	Test
LSTM dropout	79.7	77.1

A vast report on recent advances on the Penn Treebank benchmark [?], [?], [?] that shows achievement of perplexities of 70.9 on the test set.

CHAPTER 4

ANALYSIS AND DESIGN

4.1 Bayes by Backprop

Bayes by Backprop [11] is a variational inference scheme for learning the posterior distribution on the weights of a neural network. The posterior distribution on parameters of the network $\theta \in R^d$, $q(\theta)$ is typically taken to be a Gaussian with mean parameter $\mu \in R^d$ and standard deviation parameter $\sigma \in R^d$, denoted $\mathcal{N}(\theta|\mu, \sigma)$ and it's a diagonal covariance matrix. Where d is the dimensionality of the parameters of the network (usually refers to scale of millions). Let $\log p(y|\theta, x)$ be the log-likelihood of the neural network, then the network is trained by minimising the variational free energy, where $p(\theta)$ is a prior on the parameters:

$$\mathcal{L}(\theta) = E_{q(\theta)} \left[\log \frac{q(\theta)}{p(y|\theta, x)p(\theta)} \right], \quad (4.1)$$

The **algorithm 1** shows the Bayes by Backprop Monte Carlo procedure for minimising the above equation with respect to the mean and standard deviation parameters of the posterior $q(\theta)$.

Minimising the variational free energy equation is equivalent to maximising the log-likelihood $\log p(y|\theta, x)$ subject to a KL complexity term in the parameters of the network that acts as a regulariser:

$$\mathcal{L}(\theta) = -E_{q(\theta)} [\log p(y|\theta, x)] + q(\theta)p(\theta). \quad (4.2)$$

Gaussian process in a case with zero mean prior, the KL term acts as a form of weight decay on the mean parameters, where the rate of weight decay is automatically tuned by the standard deviation parameters of the prior and posterior.

Algorithm 1 Bayes by Backprop

Sample $\epsilon \sim \mathcal{N}(0, I)$, $\epsilon \in \mathbb{R}^d$.
Set network parameters to $\theta = \mu + \sigma\epsilon$.
Do forward propagation and backpropagation as normal.
Let g be the gradient with respect $L(\theta) = E_{q(\theta)} \left[\log \frac{q(\theta)}{p(y|\theta, x)p(\theta)} \right]$, to θ from backpropagation.
Let $g_\theta^{KL}, g_\mu^{KL}, g_\sigma^{KL}$ be the gradients of $\log \mathcal{N}(\theta|\mu, \sigma) - \log p(\theta)$ with respect to θ, μ and σ respectively.
Update μ according to the gradient $g + g_\theta^{KL} + g_\mu^{KL}$.
Update σ according to the gradient $(g + g_\theta^{KL})\epsilon + g_\sigma^{KL} = 0$

The uncertainty afforded by Bayes by Backprop trained networks has been used successfully for training feed-forward NN models for supervised learning and also to help exploration to reinforcement learning agents [11], [?], [?], but until now, BBB has not been applied to recurrent neural networks.

4.2 Adoption of Bayes by Backprop to RNNs

Applying Bayes by Backprop (BBB) to RNNs is illustrated in **Figure 4.1** where the weight matrices of the RNN are drawn from a distribution (learnt by BBB). However, this direct application raises two questions: when to sample the parameters of the RNN, and how to weight the contribution of the KL regulariser of (4.2). To address these concerns in the adaptation of BBB to RNNs, Algorithm 2 given where:

The variational free energy of (4.2) for an RNN on a sequence of length T is:

$$\begin{aligned} \mathcal{L}(\theta) = & -E_{q(\theta)} [\log p(y_{1:T}|\theta, x_{1:T})] \\ & + q(\theta)p(\theta), \end{aligned} \tag{4.3}$$




Figure 4.1: Adoption of Bayes by Backprop (BBB) to an RNN.

where $p(y_{1:T}|\theta, x_{1:T})$ is the likelihood of a sequence produced when the states of an unrolled RNN F_T are fed into an appropriate probability distribution. The parameters of the entire network are θ . Although the RNN is unrolled T times, each weight is penalised just once by the KL term, rather than T times. Also clear from (4.3) is that when a Monte Carlo approximation is taken to the expectation, the parameters θ should be held fixed throughout the entire sequence.

Two complications arise to the above naive derivation in practice: firstly, sequences are often long enough and models sufficiently large, that unrolling the RNN for the whole sequence is prohibitive. Secondly, to reduce variance in the gradients, more than one sequence is trained at a time. Thus the typical regime for training RNNs involves training on mini-batches of truncated sequences.

Let B be the number of mini-batches and C the number of truncated sequences (cuts), then (4.3) as:

$$\begin{aligned} \mathcal{L}(\theta) = & -E_{q(\theta)} \left[\log \prod_{b=1}^B \prod_{c=1}^C p(y^{(b,c)}|\theta, x^{(b,c)}) \right] \\ & + q(\theta)p(\theta), \end{aligned} \tag{4.4}$$

where the (b, c) superscript denotes elements of c th truncated sequence in the b th minibatch. Thus the free energy of mini-batch b of a truncated sequence c can be written as:

$$\begin{aligned}\mathcal{L}_{(b,c)}(\theta) = & -E_{q(\theta)} [\log p(y^{(b,c)} | \theta, x^{(b,c)}, s_{\text{prev}}^{(b,c)})] \\ & + w_{\text{KL}}^{(b,c)} q(\theta) p(\theta),\end{aligned}\tag{4.5}$$

where $w_{\text{KL}}^{(b,c)}$ distributes the responsibility of the KL cost among minibatches and truncated sequences (thus $\sum_{b=1}^B \sum_{c=1}^C w_{\text{KL}}^{(b,c)} = 1$), and $s_{\text{prev}}^{(b,c)}$ refers to the initial state of the RNN for the minibatch $x^{(b,c)}$. In practice, by choosing $w_{\text{KL}}^{(b,c)} = \frac{1}{CB}$ so that the KL penalty is equally distributed among all mini-batches and truncated sequences. The truncated sequences in each subsequent mini-batches are picked in order, and so $s_{\text{prev}}^{(b,c)}$ is set to the last state of the RNN for $x^{(b,c-1)}$.

Finally, the question of when to sample weights follows naturally from taking a Monte Carlo approximations to (4.5): for each minibatch, sample a fresh set of parameters.

Algorithm 2 Bayes by Backprop for RNNs

Sample $\epsilon \sim \mathcal{N}(0, I)$, $\epsilon \in \mathbb{R}^d$.

Set network parameters to $\theta = \mu + \sigma\epsilon$.

Sample a minibatch of truncated sequences (x, y) .

Do forward propagation and backpropagation as normal on minibatch.

Let g be the gradient with respect to θ from backpropagation.

Let $g_{\theta}^{KL}, g_{\mu}^{KL}, g_{\sigma}^{KL}$ be the gradients of $\log \mathcal{N}(\theta | \mu, \sigma) - \log p(\theta)$ w.r.t. θ, μ and σ respectively.

Update μ according to the gradient $\frac{g + \frac{1}{C} g_{\theta}^{KL}}{B} + \frac{g_{\mu}^{KL}}{BC}$.

Update σ according to the gradient $\left(\frac{g + \frac{1}{C} g_{\theta}^{KL}}{B} \right) \epsilon + \frac{g_{\sigma}^{KL}}{BC} = 0$

REFERENCES

1. Ghahramani, Z. Probabilistic machine learning and artificial intelligence. *Nature*, 2015. 521(7553): 452–459. ISSN 0028-0836. URL <http://dx.doi.org/10.1038/nature14541><http://10.0.4.14/nature14541>.
2. Wu, Y., Schuster, M., Chen, Z., Le, Q. V., Norouzi, M., Macherey, W., Krikun, M., Cao, Y., Gao, Q., Macherey, K., Klingner, J., Shah, A., Johnson, M., Liu, X., Kaiser, Ł., Gouws, S., Kato, Y., Kudo, T., Kazawa, H., Stevens, K., Kurian, G., Patil, N., Wang, W., Young, C., Smith, J., Riesa, J., Rudnick, A., Vinyals, O., Corrado, G., Hughes, M. and Dean, J. Google’s Neural Machine Translation System: Bridging the Gap between Human and Machine Translation. *ArXiv e-prints*, 2016: 1–23. URL <http://arxiv.org/abs/1609.08144>.
3. Amodei, D., Anubhai, R., Battenberg, E., Carl, C., Casper, J., Catanzaro, B., Chen, J., Chrzanowski, M., Coates, A., Diamos, G., Elsen, E., Engel, J., Fan, L., Fougner, C., Han, T., Hannun, A., Jun, B., LeGresley, P., Lin, L., Narang, S., Ng, A., Ozair, S., Prenger, R., Raiman, J., Satheesh, S., Seetapun, D., Sengupta, S., Wang, Y., Wang, Z., Wang, C., Xiao, B., Yogatama, D., Zhan, J. and Zhu, Z. Deep-speech 2: End-to-end speech recognition in English and Mandarin. *Jmlr W&Cp*, 2015. 48: 28. ISSN 10987576. doi:10.1145/1143844.1143891.
4. Jozefowicz, R., Vinyals, O., Schuster, M., Shazeer, N. and Wu, Y. Exploring the Limits of Language Modeling. *arXiv:1602.02410 [cs]*, 2016. URL <http://arxiv.org/abs/1602.02410>{%}5Cn<http://www.arxiv.org/pdf/1602.02410.pdf>.
5. Zaremba, W., Sutskever, I. and Vinyals, O. Recurrent Neural Network Regularization. *Iclr*, 2014. (2013): 1–8. ISSN 0157244X. doi:ng. URL <http://arxiv.org/abs/1409.2329>.
6. Lu, J., Xiong, C., Parikh, D. and Socher, R. Knowing When to Look: Adaptive Attention via A Visual Sentinel for Image Captioning. *1612.01887v1*, 2016. URL <http://arxiv.org/abs/1612.01887>.
7. Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I. and Salakhutdinov, R.

- Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine Learning Research*, 2014. 15: 1929–1958. ISSN 15337928. doi: 10.1214/12-AOS1000.
8. Gal, Y. and Ghahramani, Z. Dropout as a Bayesian Approximation : Representing Model Uncertainty in Deep Learning. *Icml*, 2015. 48: 1–10.
 9. Welling, M. and Teh, Y.-W. Bayesian Learning via Stochastic Gradient Langevin Dynamics. *Proceedings of the 28th International Conference on Machine Learning*, 2011: 681–688.
 10. Gan, Z., Li, C., Chen, C., Pu, Y., Su, Q. and Carin, L. Scalable Bayesian Learning of Recurrent Neural Networks for Language Modeling. *arXiv preprint*, 2016.
 11. Blundell, C., Cornebise, J., Kavukcuoglu, K. and Wierstra, D. Weight Uncertainty in Neural Networks. *Icml*, 2015. 37: 1613–1622. URL <http://arxiv.org/abs/1505.05424>{%}5Cn<http://www.arxiv.org/pdf/1505.05424.pdf>.
 12. Hinton, G. E., Hinton, G. E., van Camp, D. and van Camp, D. Keeping the neural networks simple by minimizing the description length of the weights. *Proceedings of the sixth annual conference on Computational learning theory - COLT '93*, 1993: 5–13. doi:10.1145/168304.168306. URL <http://portal.acm.org/citation.cfm?doid=168304.168306>.